


# Bank Vault

High-Concurrency Client-Server Network Service System

Project Plan & Technical Design

Date: December 12, 2025 (Asia/Taipei)

Team: (Server/IPC), (Protocol/Libs), (Client/Benchmarks)



**Elevator pitch:** A mini core-banking service that proves OS-level mastery: prefork multi-process server, IPC-backed shared account state with process-shared locking, a custom binary protocol with CRC integrity checks (and optional XOR payload encryption), plus a multi-threaded client load generator that reports throughput and latency percentiles.

## 1. Goals and Requirements Map

This project implements a high-concurrency client-server system with custom protocol and fault tolerance. The design intentionally maps one-to-one to the course requirements.


Requirement	Design choice in Bank Vault
Client: multi-threaded ( $\geq 100$ conns)	Load generator spawns 100+ threads with barrier start. Each thread drives a session (login + ops) and records latency.
Server: multi-process	Master preforks N worker processes (Master-Worker). Workers handle clients using epoll.
IPC required	POSIX shared memory stores account balances and global counters; process-shared locking ensures correctness.
Custom app protocol (no HTTP/WebSocket)	Binary framing: [Len][Magic][Ver][Flags][Op][Seq][CRC32][Body]. Opcodes implement banking operations.
Shared libraries	libnet (socket/epoll wrappers), libproto (encode/decode, CRC32, XOR), liblog (logging).
Security ( $\geq 1$ )	CRC32 integrity for every packet; optional XOR encryption; authentication handshake.
Reliability ( $\geq 1$ )	Heartbeat + timeouts; graceful shutdown on SIGINT; optional worker respawn on crash.

### Success criteria

- Sustain 100 concurrent client connections without crashing.
- Maintain correct balances under concurrent transfers and withdrawals (no lost updates, no negative balance).
- Demonstrate measured throughput (req/s) and latency percentiles (p50/p95/p99).
- Clean shutdown releases IPC resources (no leaked shared memory objects).

## 2. System Overview

The system consists of a multi-threaded client load generator and a prefork multi-process server. Workers run epoll loops for scalable I/O, while IPC-backed shared state provides consistent account management across processes.



### Key design decisions

- Workers share the listening socket and call accept() directly (simple prefork).
- Each worker uses epoll to handle many clients concurrently (high concurrency without one-process-per-client).
- Shared memory holds balances and counters; process-shared locking ensures data integrity.
- Protocol is length-prefixed with CRC32 for integrity and a Magic value for framing recovery.

## 3. Server Architecture (Multi-Process)

The server implements a Master-Worker architecture. The master sets up the listening socket, initializes IPC resources, then preforks workers. Workers handle I/O and business logic.

### 3.1 Process model

- Master: initializes listening socket, IPC, configuration; spawns N workers; handles SIGINT/SIGTERM; monitors worker liveness (optional respawn).
- Workers: epoll event loop; per-connection session state; protocol decode/encode; executes bank operations with shared-state locking.

### 3.2 Connection handling strategy

Recommended baseline (prefork accept): all workers inherit the same listening socket and call accept(). The OS load-balances connections across workers. This is standard, easy to implement, and sufficient for the course requirements.

Optional advanced version: master accepts connections and passes client fds to workers via a Unix domain socket using SCM\_RIGHTS (showy, but more complex).

### 3.3 Worker I/O model

- Non-blocking sockets for clients; one epoll instance per worker.
- Per-client read buffer supports partial reads and reassembly via the length-prefixed frame.
- Per-client write queue supports partial writes (EAGAIN) without blocking the worker.

### 3.4 Error handling policy

- Malformed frames: reject with error response and close connection if repeated.
- CRC failure: respond with ERR\_BAD\_CRC; optionally request retransmit; close on abuse.
- Server busy: respond with ERR\_BUSY (rate limit) or apply backpressure by pausing read events.

## 4. IPC and Shared State

IPC is used to synchronize account balances and global service statistics across worker processes. The baseline uses POSIX shared memory with a process-shared pthread mutex.

### 4.1 POSIX shared memory layout

```
// Stored in shm (mmap). All workers attach read-write.  
#define MAX_ACCOUNTS 10000  
  
typedef struct {  
    pthread_mutex_t global_lock; // PTHREAD_PROCESS_SHARED  
    uint64_t total_requests;  
    uint64_t total_errors;  
    uint32_t active_connections;  
    uint32_t shutdown_flag; // set by master on SIGINT  
    struct {  
        pthread_mutex_t lock; // per-account lock (process-shared)  
        int64_t balance_cents; // store cents to avoid floating bugs  
    } acct[MAX_ACCOUNTS];  
} vault_shm_t;
```

### 4.2 Locking rules (data consistency)

- Deposit/Withdraw: lock the account, update balance, unlock.
- Transfer: lock both accounts in a stable order (min acct id first) to prevent deadlocks.
- Counters: protected by global\_lock or atomic operations (optional).

### 4.3 Crash safety note

If a worker crashes while holding a mutex, the system can stall. As a mitigation, enable robust mutexes (pthread\_mutexattr\_setrobust) and handle EOWNERDEAD by marking shared state consistent again. This is optional but demonstrates advanced OS knowledge.

## 5. Custom Protocol Design

The protocol is a binary, length-prefixed application-layer protocol with integrity checks. It is not HTTP/WebSocket.

### 5.1 Frame format

Field	Size	Description
Length	4 bytes (uint32)	Total packet length in bytes (header + body). Network byte order.
Magic	2 bytes	0xC0DE. Used for framing recovery and quick validation.
Version	1 byte	Protocol version (start at 1).
Flags	1 byte	bit0: XOR encrypted body, bit1: error response, bit2: reserved.
OpCode	2 bytes (uint16)	Operation ID (login, deposit, withdraw, ...).
Seq	4 bytes (uint32)	Client-chosen sequence number for matching responses.
CRC32	4 bytes (uint32)	CRC32 over header (with CRC=0) + body.
Body	variable	OpCode-specific payload.

### 5.2 OpCode set

OpCode	Name	Request Body	Response Body
0x0001	LOGIN	user_len(1)+user, pass_len(1)+pass	status(2), session_id(4)
0x0002	DEPOSIT	acct_id(4), amount_cents(8)	status(2), new_balance(8)
0x0003	WITHDRAW	acct_id(4), amount_cents(8)	status(2), new_balance(8)
0x0004	TRANSFER	from(4), to(4), amount_cents(8)	status(2), from_balance(8), to_balance(8)
0x0005	BALANCE	acct_id(4)	status(2), balance(8)
0x00F0	PING	(empty)	status(2), server_time_ms(8)

### 5.3 Session key and optional XOR encryption

After LOGIN succeeds, the server returns a session\_id. A simple session key can be derived (e.g., key = CRC32(user||session\_id)) and used to XOR the Body bytes when the encrypted flag is set. This keeps the project within scope while meeting the encryption requirement.

## 6. Security and Reliability Features

The project includes an integrity check (CRC32) for every packet and implements reliability features to handle disconnections and clean shutdown.

### 6.1 Security

- Integrity: CRC32 validates each frame; packets failing CRC are rejected.
- Authentication: LOGIN handshake required before money operations.
- Encryption (optional): XOR-encrypted payload using a session-derived key (flag-controlled).

### 6.2 Reliability

- Heartbeat: clients send PING periodically; server drops idle connections after timeout.
- Timeout handling: client sets read/write timeouts and retries idempotent operations (BALANCE) when safe.
- Graceful shutdown: SIGINT triggers shutdown\_flag; workers stop accepting, finish in-flight requests, close fds, then detach and unlink shared memory.
- Worker crash isolation: a crashed worker only drops its clients; other workers continue. Optional master respawn.

## 7. Client Load Generator and Metrics

The client is a multi-threaded stress-testing tool designed to open 100+ concurrent sessions, execute randomized banking operations, and report throughput and latency statistics.

### 7.1 Threading model

- Thread-per-connection baseline: each thread maintains one TCP connection and sends sequential requests.
- Barrier start: all threads begin load at the same time to generate concurrency spikes.
- Configurable run length: operations per thread, distribution of opcodes, and target accounts.

### 7.2 Metrics output (example)

```
vault_client --host 127.0.0.1 --port 7777 --threads 100 --ops 2000 --mix  
transfer=10,withdraw=20,deposit=30,balance=40  
  
RESULT:  
total_ops=200000 ok=199872 err=128 duration=12.34s  
throughput=16206.8 req/s  
latency_ms: p50=1.2 p95=4.8 p99=9.7 max=41.3  
reconnects=3 timeouts=11
```

### 7.3 Correctness checks

- Invariant: balances never become negative (server enforces).
- Transfer atomicity: after transfer, sum of balances across involved accounts changes only by fees (if any) otherwise unchanged.
- End-to-end: compare server-side totals to client-side expected totals under controlled test cases.

## 8. Repository Layout and Build System

The project uses a Makefile-based build and shares common code via libraries to meet the modularity requirement.

### 8.1 Suggested repo layout

```
repo/
include/ # public headers shared by client + server
libnet/ # socket wrappers, epoll helpers
libproto/ # framing, CRC32, (optional) XOR
liblog/ # logging, rotating files
server/ # master-worker server, shm init, business logic
client/ # multi-thread load generator + stats
tests/ # unit/integration tests
Makefile
README.md
```

### 8.2 Make targets

- make all: builds libs + server + client
- make server: builds vault\_server
- make client: builds vault\_client
- make test: runs unit tests (protocol + CRC) and integration tests (single worker + scripted ops)
- make clean: removes build artifacts

### 8.3 Runtime commands (example)

```
# server: prefork 4 workers, 10k accounts
./vault_server --port 7777 --workers 4 --accounts 10000

# client: 100 concurrent connections, mixed ops
./vault_client --host 127.0.0.1 --port 7777 --threads 100 --ops 2000
```

## 9. Testing Plan, Milestones, and Demo Checklist

### 9.1 Testing plan

- Unit tests: protocol encode/decode, framing recovery, CRC32 correctness, XOR roundtrip.
- Concurrency tests: 100 threads, random transfers; validate invariants (no negative, atomic transfer).
- Failure tests: drop connections, delay responses, send bad CRC, send truncated frames; verify robust handling.
- Shutdown test: SIGINT during load; confirm resources cleaned (shm\_unlink) and workers exit cleanly.

### 9.2 Milestones (2-3 week schedule)

Week	Deliverables
W1	Protocol framing + libproto + skeleton server/client (connect, echo, CRC).
W2	Prefork server + epoll worker loop + shared memory + locking + core ops (deposit/withdraw/balance).
W3	Transfers + authentication + heartbeat + graceful shutdown + benchmarks + README + screenshots.

### 9.3 Team role split (example)

- Member A: server core (prefork, epoll loop, session handling, graceful shutdown).
- Member B: IPC + locking + shared memory design + correctness tests.
- Member C: protocol libs + client load generator + latency/throughput reports.

### 9.4 Demo + screenshot checklist

- Screenshot: server running with multiple worker processes (ps output).
- Screenshot: client run showing throughput and latency percentiles.
- Screenshot: CRC failure handling (client sends bad packet -> server rejects).
- Screenshot: SIGINT graceful shutdown and IPC cleanup proof (no lingering shm object).

## Appendix A. Status and Error Codes

Code	Name	Meaning
0	OK	Success
1	ERR_AUTH	Not logged in or login failed
2	ERR_BAD_CRC	Integrity check failed
3	ERR_BAD_FRAME	Malformed or unsupported packet
4	ERR_NO_ACCOUNT	Account id out of range
5	ERR_INSUFFICIENT	Not enough balance to withdraw/transfer
6	ERR_BUSY	Server overloaded or rate-limited
7	ERR_TIMEOUT	Request timed out